

# Shashaank Reddy Gurrula

Data Engineer | Azure, Databricks &  
Cloud Data Engineering

shashaankreddy01@gmail.com • +1 (309) 831-9147 • [linkedin.com/in/shashaank-reddy-g](https://www.linkedin.com/in/shashaank-reddy-g)

## PROFESSIONAL SUMMARY

Data Engineer with 5+ years of experience designing and scaling cloud-native data platforms for healthcare and financial services. Known for building robust ETL pipelines on Azure and Databricks, strengthening data governance, and delivering measurable performance gains while supporting HIPAA-sensitive analytics workloads.

## TECHNICAL SKILLS

Azure: Azure Data Factory, Azure Databricks, ADLS Gen2, Azure Synapse, Event Hubs, Stream Analytics, Logic Apps, Functions, Key Vault, DevOps • Big Data: Spark, PySpark, Delta Lake, HDFS, Hive, Kafka • Databases: Snowflake, SQL Server, Azure SQL, Oracle, Cosmos DB • Languages: SQL, PL/SQL, Python, HiveQL, Scala, Java • Modeling: Star Schema, Snowflake Schema, ER/Studio, Erwin • DevOps: Git, GitHub, Azure DevOps, CI/CD, SonarQube

## PROFESSIONAL EXPERIENCE

### Optum — Data Engineer

June 2024 – Present • Remote (San Antonio, Texas)

- Architected and scaled multi-client ingestion pipelines across Azure Data Factory, Databricks, and Snowflake to move JSON, CSV, and Parquet data into ADLS and Delta Lake with Unity Catalog RBAC for PHI-sensitive environments.
- Re-engineered the core ETL workflow with CDC and SCD Type 1/2 patterns, reducing daily processing time by 40% and enabling low-latency analytics for a risk-adjustment platform serving 17 major U.S. health plans.
- Delivered two production PySpark analytics frameworks using AI-assisted development, accelerating stakeholder reporting and tightening validation through automated QA and SonarQube compliance.
- Consolidated the platform onto Microsoft Fabric and introduced Databricks Genie to empower business users with self-service analytics on curated Gold-layer datasets.

### Capital One — Data Engineer

January 2023 – May 2024 • Remote (Normal, Illinois)

- Led a large-scale migration of tens of terabytes of data from on-prem SQL Server, MongoDB, Cassandra, and Oracle into ADLS Gen2 and Azure Synapse using ADF and Self-Hosted Integration Runtime.
- Modernized legacy Oracle ETL into cloud-native Python, PySpark, and PL/SQL pipelines with Databricks-based validation to maintain consistency across all migrated sources.
- Implemented Delta Lake and CDC/SCD patterns to improve reliability, reduce reprocessing for late-arriving data, and support near real-time reporting.
- Optimized Snowflake warehouse configuration and data modeling for enterprise BI and Power BI performance, while improving auditability through Time Travel and governance controls.

### OpenText — Data Engineer

June 2020 – December 2021 • Chennai, India

- Engineered B2B EDI ingestion pipelines for X12 850, 810, and 856 transactions, transforming purchase orders, invoices, and shipment notices into structured Parquet for downstream reconciliation.
- Automated file-arrival monitoring and pipeline triggering with Python and Azure Blob Storage SDKs, eliminating hours of manual intervention each week.
- Built Snowflake-based SCD Type 2 pipelines to track historical changes in supplier and product master data, enabling vendor and inventory analytics.
- Developed real-time event ingestion and API integration workflows that improved data freshness, compliance auditing, and partner data automation.

## EDUCATION

**Master of Science, Computer Science and Engineering — Illinois State University**

Jan 2022 – Dec 2023

**Bachelor of Technology, Computer Science Engineering (Data Analytics) — Veltech University**

Jun 2017 – May 2021

## CERTIFICATIONS

Databricks Academy Accreditation — Generative AI Fundamentals

Generative AI for Everyone — DeepLearning.AI